

Lecture#1

Sorting & Searching Arrays

CENG 102- Algorithms and Programming II,
2024-2025, Spring

Lecture 2.1

Sorting Arrays

6.8 Sorting Arrays

- Sorting data (i.e., placing the data into a particular order such as ascending or descending) is one of the most important computing applications.
- In this lecture, we will discuss what is perhaps the simplest known sorting scheme.



Performance Tip 6.4

Often, the simplest algorithms perform poorly. Their virtue is that they're easy to write, test and debug. More complex algorithms are often needed to realize maximum performance.

6.8 Sorting Arrays

- Figure 6.15 sorts the values in the elements of the 10-element array a into **ascending** order.
- The technique we use is called the **bubble sort** or the **sinking sort** because the smaller values gradually “bubble” their way upward to the top of the array like air bubbles rising in water, while the larger values sink to the bottom of the array.
- The technique is to make **several passes** through the array.
 - Several passes $\rightarrow n-1$, where the length of the array is n .

6.8 Sorting Arrays

- On each pass, successive pairs of elements (element 0 and element 1, then element 1 and element 2, etc.) are compared.
- If a pair is in increasing order (or if the values are identical), we leave the values as they are.
- If a pair is in decreasing order, their values are swapped in the array.

```
1 // Fig. 6.15: fig06_15.c
2 // Sorting an array's values into ascending order.
3 #include <stdio.h>
4 #define SIZE 10
5
6 // function main begins program execution
7 int main(void)
8 {
9     // initialize a
10    int a[SIZE] = {2, 6, 4, 8, 10, 12, 89, 68, 45, 37};
11
12    puts("Data items in original order");
13
14    // output original array
15    for (size_t i = 0; i < SIZE; ++i) {
16        printf("%4d", a[i]);
17    }
18
```

Fig. 6.15 | Sorting an array's values into ascending order. (Part I of 3.)

```
19 // bubble sort
20 // loop to control number of passes
21 for (unsigned int pass = 1; pass < SIZE; ++pass) {
22
23     // loop to control number of comparisons per pass
24     for (size_t i = 0; i < SIZE - 1; ++i) {
25
26         // compare adjacent elements and swap them if first
27         // element is greater than second element
28         if (a[i] > a[i + 1]) {
29             int hold = a[i];
30             a[i] = a[i + 1];
31             a[i + 1] = hold;
32         }
33     }
34 }
35
```

Fig. 6.15 | Sorting an array's values into ascending order. (Part 2 of 3.)


```
36     puts("\nData items in ascending order");
37
38     // output sorted array
39     for (size_t i = 0; i < SIZE; ++i) {
40         printf("%4d", a[i]);
41     }
42
43     puts("");
44 }
```

```
Data items in original order
  2  6  4  8 10 12 89 68 45 37
Data items in ascending order
  2  4  6  8 10 12 37 45 68 89
```

Fig. 6.15 | Sorting an array's values into ascending order. (Part 3 of 3.)

6.8 Sorting Arrays

- First the program compares $a[0]$ to $a[1]$, then $a[1]$ to $a[2]$, then $a[2]$ to $a[3]$, and so on until it completes the pass by comparing $a[8]$ to $a[9]$.
- Although there are 10 elements, only nine comparisons are performed.
- Because of the way the successive comparisons are made, a large value may move down the array many positions on a single pass, but a small value may move up only one position.
- On the first pass, the largest value is guaranteed to sink to the bottom element of the array, $a[9]$.

6.8 Sorting Arrays

- On the second pass, the second-largest value is guaranteed to sink to $a[8]$.
- On the ninth pass, the ninth-largest value sinks to $a[1]$.
- This leaves the smallest value in $a[0]$, so only *nine* passes of the array are needed to sort the array, even though there are *ten* elements.
- The sorting is performed by the nested for loops.

6.8 Sorting Arrays

- If a swap is necessary, it's performed by the three assignments

- `hold = a[i];`
`a[i] = a[i + 1];`
`a[i + 1] = hold;`

where the extra variable `hold` temporarily stores one of the two values being swapped.

- The swap cannot be performed with only the two assignments

- `a[i] = a[i + 1];`
`a[i + 1] = a[i];`

- If, for example, `a[i]` is 7 and `a[i + 1]` is 5, after the first assignment both values will be 5 and the value 7 will be lost—hence the need for the extra variable `hold`.

6.8 Sorting Arrays

- The chief virtue of the **bubble sort** is that **it's easy to program**.
- **However, it runs slowly** because every exchange moves an element only one position closer to its destination.
- This becomes apparent **when sorting large arrays**.
- In the exercises, we'll develop more efficient versions of the bubble sort.
- Much more efficient sorting algorithms than the bubble sort have been developed.

6.9 Case Study: Computing Mean, Median and Mode Using Arrays

- Computers are commonly used for **survey data analysis** to compile and analyze the results of surveys and opinion polls.
- Figure 6.16 uses array response initialized with 99 responses to a survey.
- Each response is a number from 1 to 9.
- The program computes the mean, median and mode of the 99 values.
- Figure 6.17 contains a sample run of this program.

```
1 // Fig. 6.16: fig06_16.c
2 // Survey data analysis with arrays:
3 // computing the mean, median and mode of the data.
4 #include <stdio.h>
5 #define SIZE 99
6
7 // function prototypes
8 void mean(const unsigned int answer[]);
9 void median(unsigned int answer[]);
10 void mode(unsigned int freq[], unsigned const int answer[]) ;
11 void bubbleSort(int a[]);
12 void printArray(unsigned const int a[]);
13
14 // function main begins program execution
15 int main(void)
16 {
17     unsigned int frequency[10] = {0}; // initialize array frequency
18
```

Fig. 6.16 | Survey data analysis with arrays: computing the mean, median and mode of the data. (Part 1 of 8.)

```
19 // initialize array response
20 unsigned int response[SIZE] =
21     {6, 7, 8, 9, 8, 7, 8, 9, 8, 9,
22      7, 8, 9, 5, 9, 8, 7, 8, 7, 8,
23      6, 7, 8, 9, 3, 9, 8, 7, 8, 7,
24      7, 8, 9, 8, 9, 8, 9, 7, 8, 9,
25      6, 7, 8, 7, 8, 7, 9, 8, 9, 2,
26      7, 8, 9, 8, 9, 8, 9, 7, 5, 3,
27      5, 6, 7, 2, 5, 3, 9, 4, 6, 4,
28      7, 8, 9, 6, 8, 7, 8, 9, 7, 8,
29      7, 4, 4, 2, 5, 3, 8, 7, 5, 6,
30      4, 5, 6, 1, 6, 5, 7, 8, 7};
31
32 // process responses
33 mean(response);
34 median(response);
35 mode(frequency, response);
36 }
37
```

Fig. 6.16 | Survey data analysis with arrays: computing the mean, median and mode of the data. (Part 2 of 8.)

```
38 // calculate average of all response values
39 void mean(const unsigned int answer[])
40 {
41     printf("%s\n%s\n%s\n", "*****", "  Mean", "*****");
42
43     unsigned int total = 0; // variable to hold sum of array elements
44
45     // total response values
46     for (size_t j = 0; j < SIZE; ++j) {
47         total += answer[j];
48     }
49
50     printf("The mean is the average value of the data\n"
51           "items. The mean is equal to the total of\n"
52           "all the data items divided by the number\n"
53           "of data items (%u). The mean value for\n"
54           "this run is: %u / %u = %.4f\n\n",
55           SIZE, total, SIZE, (double) total / SIZE);
56 }
57
```

Fig. 6.16 | Survey data analysis with arrays: computing the mean, median and mode of the data. (Part 3 of 8.)

```
58 // sort array and determine median element's value
59 void median(unsigned int answer[])
60 {
61     printf("\n%s\n%s\n%s\n%s",
62           "*****", " Median", "*****",
63           "The unsorted array of responses is");
64
65     printArray(answer); // output unsorted array
66
67     bubbleSort(answer); // sort array
68
69     printf("%s", "\n\nThe sorted array is");
70     printArray(answer); // output sorted array
71
72     // display median element
73     printf("\n\nThe median is element %u of\n"
74           "the sorted %u element array.\n"
75           "For this run the median is %u\n\n",
76           SIZE / 2, SIZE, answer[SIZE / 2]);
77 }
78
```

Fig. 6.16 | Survey data analysis with arrays: computing the mean, median and mode of the data. (Part 4 of 8.)

```

79 // determine most frequent response
80 void mode(unsigned int freq[], const unsigned int answer[])
81 {
82     printf("\n%s\n%s\n%s\n", "*****", "   Mode", "*****");
83
84     // initialize frequencies to 0
85     for (size_t rating = 1; rating <= 9; ++rating) {
86         freq[rating] = 0;
87     }
88
89     // summarize frequencies
90     for (size_t j = 0; j < SIZE; ++j) {
91         ++freq[answer[j]];
92     }
93
94     // output headers for result columns
95     printf("%s%11s%19s\n\n%54s\n%54s\n\n",
96           "Response", "Frequency", "Histogram",
97           "1    1    2    2", "5    0    5    0    5");
98
99     // output results
100    unsigned int largest = 0; // represents largest frequency
101    unsigned int modeValue = 0; // represents most frequent response

```

Fig. 6.16 | Survey data analysis with arrays: computing the mean, median and mode of the data. (Part 5 of 8.)

```
102
103     for (rating = 1; rating <= 9; ++rating) {
104         printf("%8u%11u", rating, freq[rating]);
105
106         // keep track of mode value and largest frequency value
107         if (freq[rating] > largest) {
108             largest = freq[rating];
109             modeValue = rating;
110         }
111
112         // output histogram bar representing frequency value
113         for (unsigned int h = 1; h <= freq[rating]; ++h) {
114             printf("%s", "*");
115         }
116
117         puts(""); // being new line of output
118     }
119
120     // display the mode value
121     printf("\nThe mode is the most frequent value.\n"
122           "For this run the mode is %u which occurred"
123           " %u times.\n", modeValue, largest);
124 }
```

Fig. 6.16 | Survey data analysis with arrays: computing the mean, median and mode of the data. (Part 6 of 8.)

```
I25
I26 // function that sorts an array with bubble sort algorithm
I27 void bubbleSort(unsigned int a[])
I28 {
I29     // loop to control number of passes
I30     for (unsigned int pass = 1; pass < SIZE; ++pass) {
I31
I32         // loop to control number of comparisons per pass
I33         for (size_t j = 0; j < SIZE - 1; ++j) {
I34
I35             // swap elements if out of order
I36             if (a[j] > a[j + 1]) {
I37                 unsigned int hold = a[j];
I38                 a[j] = a[j + 1];
I39                 a[j + 1] = hold;
I40             }
I41         }
I42     }
I43 }
I44
```

Fig. 6.16 | Survey data analysis with arrays: computing the mean, median and mode of the data. (Part 7 of 8.)

```
145 // output array contents (20 values per row)
146 void printArray(const unsigned int a[])
147 {
148     // output array contents
149     for (size_t j = 0; j < SIZE; ++j) {
150
151         if (j % 20 == 0) { // begin new line every 20 values
152             puts("");
153         }
154
155         printf("%2u", a[j]);
156     }
157 }
```

Fig. 6.16 | Survey data analysis with arrays: computing the mean, median and mode of the data. (Part 8 of 8.)

```
*****
```

```
Mean
```

```
*****
```

```
The mean is the average value of the data  
items. The mean is equal to the total of  
all the data items divided by the number  
of data items (99). The mean value for  
this run is: 681 / 99 = 6.8788
```

Fig. 6.17 | Sample run for the survey data analysis program. (Part 1 of 3.)

```
*****
Median
*****
The unsorted array of responses is
6 7 8 9 8 7 8 9 8 9 7 8 9 5 9 8 7 8 7 8
6 7 8 9 3 9 8 7 8 7 7 8 9 8 9 8 9 7 8 9
6 7 8 7 8 7 9 8 9 2 7 8 9 8 9 8 9 7 5 3
5 6 7 2 5 3 9 4 6 4 7 8 9 6 8 7 8 9 7 8
7 4 4 2 5 3 8 7 5 6 4 5 6 1 6 5 7 8 7

The sorted array is
1 2 2 2 3 3 3 3 4 4 4 4 4 5 5 5 5 5 5 5
5 6 6 6 6 6 6 6 6 6 7 7 7 7 7 7 7 7 7 7
7 7 7 7 7 7 7 7 7 7 7 7 7 8 8 8 8 8 8 8
8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9

The median is element 49 of
the sorted 99 element array.
For this run the median is 7
```

Fig. 6.17 | Sample run for the survey data analysis program. (Part 2 of 3.)


```

*****
  Mode
*****
Response  Frequency      Histogram

                    5      1      1      2      2
                        0      5      0      5

1           1           *
2           3          ***
3           4         ****
4           5         *****
5           8         ********
6           9         **********
7          23        *****************
8          27        *****************
9          19        *************

```

The mode is the most frequent value.
 For this run the mode is 8 which occurred 27 times.

Fig. 6.17 | Sample run for the survey data analysis program. (Part 3 of 3.)

6.9 Case Study: Computing Mean, Median and Mode Using Arrays

Mean

- The *mean* is the arithmetic average of the 99 values.
- Function `mean` (Fig. 6.16) computes the mean by totaling the 99 elements and dividing the result by 99.

6.9 Case Study: Computing Mean, Median and Mode Using Arrays (Cont.)

Median

- The median is the “*middle* value.”
- Function `median` determines the median by calling function `bubbleSort` to sort the array of responses into ascending order, then picking `answer[SIZE / 2]` (the middle element) of the sorted array.
- When the number of elements is even, the median should be calculated as the mean of the two middle elements.
- Function `median` does not currently provide this capability.

6.9 Case Study: Computing Mean, Median and Mode Using Arrays (Cont.)

Mode

- The *mode* is the *value that occurs most frequently* among the 99 responses.
- Function `mode` determines the mode by counting the number of responses of each type, then selecting the value with the greatest count.
- This version of function `mode` does not handle a tie.
- Function `mode` also produces a histogram to aid in determining the mode graphically.

Lecture 2.2

Searching Arrays

6.10 Searching Arrays

- The process of finding a particular element of an array is called **searching**.
- In this section we discuss two searching techniques—the simple **linear search** technique and the more efficient (but more complex) **binary search** technique.

6.10 Searching Arrays (Cont.)

Searching an Array with Linear Search

- The linear search (Fig. 6.18) compares each element of the array with the **search key**.
- Because the array is not in any particular order, it's just as likely that the value will be found in the first element as in the last.
- On average, therefore, the program will have to compare the search key with *half* the elements of the array.

```
1  // Fig. 6.18: fig06_18.c
2  // Linear search of an array.
3  #include <stdio.h>
4  #define SIZE 100
5
6  // function prototype
7  size_t linearSearch(const int array[], int key, size_t size);
8
9  // function main begins program execution
10 int main(void)
11 {
12     int a[SIZE]; // create array a
13
14     // create some data
15     for (size_t x = 0; x < SIZE; ++x) {
16         a[x] = 2 * x;
17     }
18
19     printf("Enter integer search key: ");
20     int searchKey; // value to locate in array a
21     scanf("%d", &searchKey);
22
```

Fig. 6.18 | Linear search of an array. (Part I of 3.)

```
23     // attempt to locate searchKey in array a
24     size_t index = linearSearch(a, searchKey, SIZE);
25
26     // display results
27     if (index != -1) {
28         printf("Found value at index %d\n", index);
29     }
30     else {
31         puts("Value not found");
32     }
33 }
34
```

Fig. 6.18 | Linear search of an array. (Part 2 of 3.)

```
35 // compare key to every element of array until the location is found
36 // or until the end of array is reached; return index of element
37 // if key is found or -1 if key is not found
38 size_t linearSearch(const int array[], int key, size_t size)
39 {
40     // loop through array
41     for (size_t n = 0; n < size; ++n) {
42         if (array[n] == key) {
43             return n; // return location of key
44         }
45     }
46     return -1; // key not found
47 }
48
49
```

Enter integer search key: 36
Found value at index 18

Enter integer search key: 37
Value not found

Fig. 6.18 | Linear search of an array. (Part 3 of 3.)

6.10 Searching Arrays (Cont.)

Searching an Array with Binary Search

- The **linear searching** method **works well for *small or unsorted* arrays**.
- However, for large arrays linear searching is *inefficient*.
- **If the array is sorted**, the high-speed **binary search technique can be used**.
- The binary search algorithm **eliminates** from consideration ***one-half*** of the elements in a sorted array after each comparison.
- The algorithm locates the ***middle*** element of the array and **compares** it to the search **key**.

6.10 Searching Arrays (Cont.)

- If they're equal, the search key is found and the index of that element is returned.
- If they're not equal, the problem is reduced to searching *one-half* of the array.
- If the search key is less than the middle element of the array, the *first half* of the array is searched, otherwise the *second half* of the array is searched.
- If the search key is not found in the specified subarray (piece of the original array), the algorithm is repeated on one-quarter of the original array.

6.10 Searching Arrays (Cont.)

- The search continues until the search key is equal to the middle element of a subarray, or until the subarray consists of one element that is not equal to the search key (i.e., the search key is not found).
- In a worst case-scenario, searching an array of 1023 elements takes *only* 10 comparisons using a binary search.
- Repeatedly dividing 1024 by 2 yields the values 512, 256, 128, 64, 32, 16, 8, 4, 2 and 1.
- The number 1024 (2^{10}) is divided by 2 only 10 times to get the value 1.

6.10 Searching Arrays (Cont.)

- An array of 1048576 (2^{20}) elements takes a maximum of *only* 20 comparisons to find the search key.
- An array of one billion elements takes a maximum of *only* 30 comparisons to find the search key.
- This is a tremendous increase in performance over the linear search that required comparing the search key to an average of half of the array elements.
- For a one-billion-element array, this is a difference between an average of 500 million comparisons and a maximum of 30 comparisons!

6.10 Searching Arrays (Cont.)

- The maximum comparisons for any array can be determined by finding the first power of 2 greater than the number of array elements.
- Figure 6.19 presents the *iterative* version of function `binarySearch`
- The function receives four arguments—an integer array `b` to be searched, an integer `searchKey`, the `low` array index and the `high` array index (these define the portion of the array to be searched).
- If the search key does *not* match the middle element of a subarray, the `low` index or `high` index is modified so that a smaller subarray can be searched.

6.10 Searching Arrays (Cont.)

- If the search key is *less than* the middle element, the high index is set to `middle - 1` and the search is continued on the elements from `low` to `middle - 1`.
- If the search key is *greater than* the middle element, the low index is set to `middle + 1` and the search is continued on the elements from `middle + 1` to `high`.
- The program uses an array of 15 elements.
- The first power of 2 greater than the number of elements in this array is 16 (2^4), so no more than 4 comparisons are required to find the search key.

6.10 Searching Arrays (Cont.)

- The program uses function `printHeader` to output the array indices and function `printRow` to output each subarray during the binary search process.
- The middle element in each subarray is marked with an asterisk (*) to indicate the element to which the search key is compared.

```
1 // Fig. 6.19: fig06_19.c
2 // Binary search of a sorted array.
3 #include <stdio.h>
4 #define SIZE 15
5
6 // function prototypes
7 size_t binarySearch(const int b[], int searchKey, size_t low, size_t high);
8 void printHeader(void);
9 void printRow(const int b[], size_t low, size_t mid, size_t high);
10
11 // function main begins program execution
12 int main(void)
13 {
14     int a[SIZE]; // create array a
15
16     // create data
17     for (size_t i = 0; i < SIZE; ++i) {
18         a[i] = 2 * i;
19     }
20
21     printf("%s", "Enter a number between 0 and 28: ");
22     int key; // value to locate in array a
23     scanf("%d", &key);
24
```

Fig. 6.19 | Binary search of a sorted array. (Part I of 7.)

```
25     printHeader();
26
27     // search for key in array a
28     size_t result = binarySearch(a, key, 0, SIZE - 1);
29
30     // display results
31     if (result != -1) {
32         printf("\n%d found at index %d\n", key, result);
33     }
34     else {
35         printf("\n%d not found\n", key);
36     }
37 }
38
39 // function to perform binary search of an array
40 size_t binarySearch(const int b[], int searchKey, size_t low, size_t high)
41 {
42     // loop until low index is greater than high index
43     while (low <= high) {
44
45         // determine middle element of subarray being searched
46         size_t middle = (low + high) / 2;
47
```

Fig. 6.19 | Binary search of a sorted array. (Part 2 of 7.)

```
48     // display subarray used in this loop iteration
49     printRow(b, low, middle, high);
50
51     // if searchKey matched middle element, return middle
52     if (searchKey == b[middle]) {
53         return middle;
54     }
55
56     // if searchKey is less than middle element, set new high
57     else if (searchKey < b[middle]) {
58         high = middle - 1; // search low end of array
59     }
60
61     // if searchKey is greater than middle element, set new low
62     else {
63         low = middle + 1; // search high end of array
64     }
65 } // end while
66
67 return -1; // searchKey not found
68 }
69
```

Fig. 6.19 | Binary search of a sorted array. (Part 3 of 7.)

```
70 // Print a header for the output
71 void printHeader(void)
72 {
73     puts("\nIndices:");
74
75     // output column head
76     for (unsigned int i = 0; i < SIZE; ++i) {
77         printf("%3u ", i);
78     }
79
80     puts(""); // start new line of output
81
82     // output line of - characters
83     for (unsigned int i = 1; i <= 4 * SIZE; ++i) {
84         printf("%s", "-");
85     }
86
87     puts(""); // start new line of output
88 }
89
```

Fig. 6.19 | Binary search of a sorted array. (Part 4 of 7.)

```
90 // Print one row of output showing the current
91 // part of the array being processed.
92 void printRow(const int b[], size_t low, size_t mid, size_t high)
93 {
94     // loop through entire array
95     for (size_t i = 0; i < SIZE; ++i) {
96
97         // display spaces if outside current subarray range
98         if (i < low || i > high) {
99             printf("%s", "    ");
100        }
101        else if (i == mid) { // display middle element
102            printf("%3d*", b[i]); // mark middle value
103        }
104        else { // display other elements in subarray
105            printf("%3d ", b[i]);
106        }
107    }
108
109    puts(""); // start new line of output
110 }
```

Fig. 6.19 | Binary search of a sorted array. (Part 5 of 7.)

Enter a number between 0 and 28: 25

Indices:

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
<hr/>														
0	2	4	6	8	10	12	14*	16	18	20	22	24	26	28
								16	18	20	22*	24	26	28
												24	26*	28
												24*		

25 not found

Fig. 6.19 | Binary search of a sorted array. (Part 6 of 7.)

Enter a number between 0 and 28: **8**

Indices:

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
<hr/>														
0	2	4	6	8	10	12	14*	16	18	20	22	24	26	28
0	2	4	6*	8	10	12								
				8	10*	12								
				8*										

8 found at index 4

Enter a number between 0 and 28: **6**

Indices:

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
<hr/>														
0	2	4	6	8	10	12	14*	16	18	20	22	24	26	28
0	2	4	6*	8	10	12								

6 found at index 3

Fig. 6.19 | Binary search of a sorted array. (Part 7 of 7.)